

nhanes_bp

February 24, 2020

1 Multilevel analysis of NHANES blood pressure data

In this notebook, we fit mixed models to blood pressure data from NHANES. The study has data for two blood pressure measurement types (systolic BP and diastolic BP), with up to 4 repeated measures per subject for each type.

The data files can be obtained from these links:

https://wwwn.cdc.gov/Nchs/Nhanes/2011-2012/DEMO_G.XPT

https://wwwn.cdc.gov/Nchs/Nhanes/2011-2012/BPX_G.XPT https://wwwn.cdc.gov/Nchs/Nhanes/2011-2012/BMX_G.XPT

```
[1]: import statsmodels.api as sm
import pandas as pd
import numpy as np
```

```
/nfs/kshedden/python3/lib/python3.7/site-
packages/statsmodels/compat/pandas.py:23: FutureWarning: The Panel class is
removed from pandas. Accessing it from the top-level namespace will also be
removed in the next version
```

```
data_klasses = (pandas.Series, pandas.DataFrame, pandas.Panel)
```

First, load and merge the data sets. These are SAS Xport format files, which can be read with Pandas.

```
[2]: demog = pd.read_sas("../data/DEMO_G.XPT")
bpx = pd.read_sas("../data/BPX_G.XPT")
bmx = pd.read_sas("../data/BMX_G.XPT")
df = pd.merge(demog, bpx, left_on="SEQN", right_on="SEQN")
df = pd.merge(df, bmx, left_on="SEQN", right_on="SEQN")
```

Next we convert the data from wide to long, pivoting the four BP measures from columns to rows.

```
[3]: syvars = ["BPXSY%d" % j for j in (1,2,3,4)]
divars = ["BPXDI%d" % j for j in (1,2,3,4)]
vvars = syvars + divars
idvars = ['SEQN', 'RIDAGEYR', 'RIAGENDR', 'BMXBMI']
dx = pd.melt(df, id_vars=idvars, value_vars=vvars,
             var_name='bpvar', value_name='bp')
```

We drop rows where any of the variables are missing. Multilevel modeling can accommodate missing data, but here we use a basic complete case analysis.

```
[4]: dx = dx.sort_values(by='SEQN')
dx = dx.reset_index(drop=True)
dx['SEQN'] = dx.SEQN.astype(np.int)
dx = dx.dropna()
```

Since we have pivoted all BP measures to rows, we will need variables telling us whether we are looking at systolic (SY) or diastolic (DI) blood pressure, and we need a way to know the order of the BP values within each person. These repeated measures are not exchangeable, since the BP readings tend to drop slightly as people relax.

```
[5]: # Blood pressure type (systolic or diastolic)
dx["bpt"] = dx.bpvar.str[3:5]

dx["bpi"] = dx.bpvar.str[5].astype(np.int)
dx["female"] = (dx.RIAGENDR == 2).astype(np.int)

di_mean = dx.loc[dx.bpt=="DI", :].groupby("SEQN")["bp"].aggregate(np.mean)
di_mean.name = "di_mean"
dx = pd.merge(dx, di_mean, left_on="SEQN", right_index=True)

print(dx.head())
```

	SEQN	RIDAGEYR	RIAGENDR	BMXBMI	bpvar	bp	bpt	bpi	female	\
0	62161	22.0	1.0	23.3	BPXSY1	110.0	SY	1	0	
2	62161	22.0	1.0	23.3	BPXDI2	68.0	DI	2	0	
3	62161	22.0	1.0	23.3	BPXSY3	118.0	SY	3	0	
4	62161	22.0	1.0	23.3	BPXDI3	74.0	DI	3	0	
5	62161	22.0	1.0	23.3	BPXSY2	104.0	SY	2	0	

	di_mean
0	74.666667
2	74.666667
3	74.666667
4	74.666667
5	74.666667

We subsample the data to make the script run faster. Statsmodels MixedLM is unfortunately not very fast.

```
[6]: dx = dx.iloc[0:5000, :]
```

Fit a linear mean structure model using OLS. The variance structure of this model is misspecified. Since this is a linear model, the coefficient point estimates are still meaningful despite the variance mis-specification.

```
[7]: model1 = sm.OLS.from_formula("bp ~ RIDAGEYR + female + C(bpt) + BMXBMI", dx)
result1 = model1.fit()
print(result1.summary())
```

OLS Regression Results

=====

```

Dep. Variable:          bp    R-squared:          0.762
Model:                  OLS   Adj. R-squared:       0.762
Method:                 Least Squares   F-statistic:         4005.
Date:                  Mon, 24 Feb 2020   Prob (F-statistic):    0.00
Time:                  15:00:20   Log-Likelihood:       -20686.
No. Observations:      5000   AIC:                  4.138e+04
Df Residuals:          4995   BIC:                  4.141e+04
Df Model:              4
Covariance Type:       nonrobust

```

```

=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----
Intercept      46.1460      0.863      53.463      0.000      44.454      47.838
C(bpt)[T.SY]   52.0088      0.429     121.281      0.000      51.168      52.849
RIDAGEYR        0.2911      0.010      27.819      0.000        0.271        0.312
female        -2.1229      0.434      -4.888      0.000      -2.974      -1.272
BMXBMI         0.3696      0.032      11.584      0.000        0.307        0.432
=====
Omnibus:              424.509   Durbin-Watson:          1.178
Prob(Omnibus):         0.000   Jarque-Bera (JB):       2687.614
Skew:                 -0.060   Prob(JB):               0.00
Kurtosis:             6.590   Cond. No.               209.
=====

```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Fit a mixed model to the systolic data with a simple random intercept per subject. Then calculate the ICC.

```

[8]: ds2 = dx.loc[dx.bpt == "SY"]
      model2 = sm.MixedLM.from_formula("bp ~ RIDAGEYR + female + BMXBMI",
                                      groups="SEQN", data=ds2)
      result2 = model2.fit()
      icc2 = result2.cov_re / (result2.cov_re + result2.scale)
      print(result2.summary())
      print("icc=%f\n" % icc2.values.flat[0])

```

Mixed Linear Model Regression Results

```

=====
Model:              MixedLM Dependent Variable: bp
No. Observations:  2500   Method:              REML
No. Groups:        837   Scale:               16.2140
Min. group size:    1    Likelihood:          -8563.6384
Max. group size:    3    Converged:           Yes
Mean group size:    3.0
=====
              Coef.  Std.Err.   z      P>|z| [0.025  0.975]
-----

```

```
-----
Intercept      96.233      1.960 49.100 0.000 92.391 100.074
RIDAGEYR        0.406      0.025 16.529 0.000  0.358  0.454
female         -2.765      1.018 -2.715 0.007 -4.762 -0.769
BMXBMI          0.289      0.075  3.864 0.000  0.143  0.436
SEQN Var       206.352      3.158
=====
```

icc=0.927150

Partial out the mean diastolic blood pressure per subject. This leads to slightly weaker random effects.

```
[9]: model3 = sm.MixedLM.from_formula("bp ~ RIDAGEYR + female + BMXBMI + di_mean",
                                     groups="SEQN", data=ds2)
result3 = model3.fit()
icc3 = result3.cov_re / (result3.cov_re + result3.scale)
print(result3.summary())
print("icc=%f\n" % icc3.values.flat[0])
```

Mixed Linear Model Regression Results

```
=====
Model:              MixedLM Dependent Variable: bp
No. Observations: 2500   Method:              REML
No. Groups:          837   Scale:              16.2135
Min. group size:     1     Likelihood:         -8517.7505
Max. group size:     3     Converged:           Yes
Mean group size:     3.0
=====
```

```
-----
              Coef.  Std.Err.   z    P>|z| [0.025 0.975]
-----
Intercept      79.975      2.450 32.639 0.000 75.173 84.778
RIDAGEYR        0.345      0.024 14.443 0.000  0.299  0.392
female         -2.238      0.963 -2.324 0.020 -4.126 -0.351
BMXBMI          0.140      0.072  1.932 0.053 -0.002  0.281
di_mean         0.337      0.033 10.122 0.000  0.272  0.402
SEQN Var       183.365      2.816
=====
```

icc=0.918761

Fit a mixed model to the diastolic data with a random intercept per subject. Then calculate the ICC.

```
[10]: ds3 = dx.loc[dx.bpt == "DI"]
model4 = sm.MixedLM.from_formula("bp ~ RIDAGEYR + female + BMXBMI",
                                  groups="SEQN", data=ds3)
result4 = model4.fit()
```

```
icc4 = result4.cov_re / (result4.cov_re + result4.scale)
print(result4.summary())
print("icc=%f\n" % icc4.values.flat[0])
```

Mixed Linear Model Regression Results

```
=====
Model:                MixedLM Dependent Variable: bp
No. Observations: 2500   Method:                REML
No. Groups:           837   Scale:                36.7592
Min. group size:      1    Likelihood:            -9229.4228
Max. group size:      3    Converged:              Yes
Mean group size:      3.0

-----
                Coef.   Std.Err.    z    P>|z|  [0.025  0.975]
-----
Intercept      48.224    1.926  25.045  0.000  44.450  51.998
RIDAGEYR        0.178    0.024   7.401  0.000   0.131   0.226
female         -1.562    1.001  -1.561  0.119  -3.523   0.399
BMXBMI          0.445    0.074   6.044  0.000   0.300   0.589
SEQN Var      192.039    2.024

=====
```

```
icc=0.839337
```

Fit a mixed model to the diastolic data with a random intercept per subject.

```
[11]: ds3 = dx.loc[dx.bpt == "DI"]
model5 = sm.MixedLM.from_formula("bp ~ RIDAGEYR + female + BMXBMI + bpi",
                                groups="SEQN", re_formula="1+bpi",
                                data=ds3)

result5 = model5.fit()
print(result5.summary())
```

```
/nfs/kshedden/python3/lib/python3.7/site-packages/statsmodels/base/model.py:512:
ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check
mle_retvals
  "Check mle_retvals", ConvergenceWarning)
/nfs/kshedden/python3/lib/python3.7/site-
packages/statsmodels/regression/mixed_linear_model.py:2059: ConvergenceWarning:
Retrying MixedLM optimization with lbfgs
  ConvergenceWarning)
/nfs/kshedden/python3/lib/python3.7/site-packages/statsmodels/base/model.py:512:
ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check
mle_retvals
  "Check mle_retvals", ConvergenceWarning)
/nfs/kshedden/python3/lib/python3.7/site-
packages/statsmodels/regression/mixed_linear_model.py:2059: ConvergenceWarning:
```

```

Retrying MixedLM optimization with cg
ConvergenceWarning)
/nfs/kshedden/python3/lib/python3.7/site-packages/statsmodels/base/model.py:512:
ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check
mle_retvals
"Check mle_retvals", ConvergenceWarning)
/nfs/kshedden/python3/lib/python3.7/site-
packages/statsmodels/regression/mixed_linear_model.py:2063: ConvergenceWarning:
MixedLM optimization failed, trying a different optimizer may help.
warnings.warn(msg, ConvergenceWarning)
/nfs/kshedden/python3/lib/python3.7/site-
packages/statsmodels/regression/mixed_linear_model.py:2075: ConvergenceWarning:
Gradient optimization failed, |grad| = 11.616035
warnings.warn(msg, ConvergenceWarning)

```

Mixed Linear Model Regression Results

```

=====
Model:                MixedLM Dependent Variable: bp
No. Observations:    2500    Method:                REML
No. Groups:          837     Scale:                35.6655
Min. group size:     1       Likelihood:           -9204.0121
Max. group size:     3       Converged:             No
Mean group size:     3.0
-----

```

	Coef.	Std.Err.	z	P> z	[0.025	0.975]
Intercept	49.150	1.865	26.359	0.000	45.496	52.805
RIDAGEYR	0.193	0.023	8.253	0.000	0.147	0.239
female	-1.512	0.965	-1.566	0.117	-3.404	0.380
BMXBMI	0.422	0.071	5.942	0.000	0.283	0.561
bpi	-0.448	0.149	-3.010	0.003	-0.739	-0.156
SEQN Var	136.382	2.444				
SEQN x bpi Cov	11.527	0.406				
bpi Var	1.244	0.199				

```

=====

```

Fit the same model as above, now centering the bpi (index) variable.

```

[12]: ds3.loc[:, "bpi_cen"] = ds3.loc[:, "bpi"] - 1
model6 = sm.MixedLM.from_formula("bp ~ RIDAGEYR + female + BMXBMI + bpi",
                                groups="SEQN", re_formula="1+bpi_cen",
                                data=ds3)

result6 = model6.fit()
print(result6.summary())

```

```

/nfs/kshedden/python3/lib/python3.7/site-packages/pandas/core/indexing.py:376:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.

```

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
self.obj[key] = _infer_fill_value(value)
/nfs/kshedden/python3/lib/python3.7/site-packages/pandas/core/indexing.py:494:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
self.obj[item] = s
```

Mixed Linear Model Regression Results

```
=====
Model:                MixedLM    Dependent Variable:    bp
No. Observations:    2500        Method:                REML
No. Groups:          837         Scale:                31.6095
Min. group size:     1           Likelihood:            -9199.2717
Max. group size:     3           Converged:              Yes
Mean group size:     3.0

-----
                Coef.  Std.Err.    z    P>|z|  [0.025  0.975]
-----
Intercept          49.156    1.891  25.997  0.000   45.450   52.861
RIDAGEYR            0.191    0.024   8.077  0.000    0.145    0.237
female             -1.524    0.977  -1.559  0.119   -3.439    0.392
BMXBMI              0.425    0.072   5.909  0.000    0.284    0.566
bpi                -0.458    0.156  -2.946  0.003   -0.763   -0.153
SEQN Var          170.158    2.448
SEQN x bpi_cen Cov  8.576    0.399
bpi_cen Var        4.835    0.253
=====
```

Fit a mixed model to both types of BP jointly with a random intercept per subject. Note that the random intercept is shared between the two types of blood pressure (systolic and diastolic).

```
[13]: model7 = sm.MixedLM.from_formula("bp ~ RIDAGEYR + female + C(bpt) + BMXBMI",
                                     groups="SEQN", data=dx)
result7 = model7.fit()
print(result7.summary())
```

Mixed Linear Model Regression Results

```
=====
Model:                MixedLM    Dependent Variable:    bp
No. Observations:    5000        Method:                REML
No. Groups:          837         Scale:                110.9918
```

Min. group size: 2 Likelihood: -19709.8882
 Max. group size: 6 Converged: Yes
 Mean group size: 6.0

	Coef.	Std.Err.	z	P> z	[0.025	0.975]
Intercept	46.215	1.592	29.038	0.000	43.096	49.335
C(bpt)[T.SY]	52.009	0.298	174.537	0.000	51.425	52.593
RIDAGEYR	0.292	0.020	14.715	0.000	0.253	0.331
female	-2.157	0.823	-2.620	0.009	-3.771	-0.544
BMXBMI	0.367	0.061	6.068	0.000	0.249	0.486
SEQN Var	119.727	0.705				

Fit a mixed model to both types of BP with a subject random intercept and a unique random effect per BP type with common variance.

```
[14]: model8 = sm.MixedLM.from_formula("bp ~ RIDAGEYR + female + C(bpt) + BMXBMI",
                                     groups="SEQN", re_formula="1",
                                     vc_formula={"bpt": "0+C(bpt)"},
                                     data=dx)

result8 = model8.fit()
print(result8.summary())
```

Mixed Linear Model Regression Results

Model: MixedLM Dependent Variable: bp
 No. Observations: 5000 Method: REML
 No. Groups: 837 Scale: 26.4863
 Min. group size: 2 Likelihood: -17909.6150
 Max. group size: 6 Converged: Yes
 Mean group size: 6.0

	Coef.	Std.Err.	z	P> z	[0.025	0.975]
Intercept	46.221	1.613	28.656	0.000	43.060	49.383
C(bpt)[T.SY]	52.018	0.598	86.919	0.000	50.845	53.191
RIDAGEYR	0.292	0.020	14.716	0.000	0.253	0.331
female	-2.164	0.824	-2.628	0.009	-3.779	-0.550
BMXBMI	0.367	0.061	6.059	0.000	0.248	0.486
SEQN Var	63.569	1.530				
bpt Var	140.980	1.593				

Fit a mixed model to both types of BP with a subject random intercept and a unique random effect per BP type with unique variance.


```
[15]: dx["sy"] = (dx.bpt == "SY").astype(np.int)
dx["di"] = (dx.bpt == "DI").astype(np.int)
model9 = sm.MixedLM.from_formula("bp ~ RIDAGEYR + female + C(bpt) + BMXBMI",
                                groups="SEQN", re_formula="1",
                                vc_formula={"sy": "0+sy", "di": "0+di"},
                                data=dx)

result9 = model9.fit()
print(result9.summary())
```

Mixed Linear Model Regression Results

```
=====
Model:                MixedLM Dependent Variable: bp
No. Observations: 5000  Method:                REML
No. Groups:          837  Scale:                26.4867
Min. group size:     2    Likelihood:          -17909.4687
Max. group size:     6    Converged:          Yes
Mean group size:     6.0

-----
```

	Coef.	Std.Err.	z	P> z	[0.025	0.975]
Intercept	46.274	1.614	28.669	0.000	43.111	49.438
C(bpt) [T.SY]	52.018	0.598	86.914	0.000	50.845	53.191
RIDAGEYR	0.289	0.021	14.030	0.000	0.249	0.329
female	-2.148	0.824	-2.608	0.009	-3.763	-0.534
BMXBMI	0.369	0.061	6.083	0.000	0.250	0.488
SEQN Var	63.540	1.529				
di Var	137.038	2.105				
sy Var	144.954	2.168				

```
=====
```

Below we consider the possibility that there may be heteroscedasticity between the two blood pressure types. That is, systolic blood pressure measurements may be more variable than diastolic measurements, or vice versa. This analysis is a bit awkward to conduct. Below we fit two models, one in which diastolic measurements are allowed to be more variable than systolic measurements, and one in which systolic measurements are allowed to be more variable than diastolic measurements. In theory, a variance parameters should be equal to zero in one of these models, revealing which type of blood pressure has more variability.

```
[16]: dx["sy1"] = (dx.bpvar == "BPXSY1").astype(np.int)
dx["sy2"] = (dx.bpvar == "BPXSY2").astype(np.int)
dx["sy3"] = (dx.bpvar == "BPXSY3").astype(np.int)
dx["di1"] = (dx.bpvar == "BPXDI1").astype(np.int)
dx["di2"] = (dx.bpvar == "BPXDI2").astype(np.int)
dx["di3"] = (dx.bpvar == "BPXDI3").astype(np.int)
model10 = sm.MixedLM.from_formula("bp ~ RIDAGEYR + female + C(bpt) + BMXBMI",
                                   groups="SEQN", re_formula="1",
                                   vc_formula={"sy": "0+sy", "di": "0+di",
```

```

                                "dye": "0+di1+di2+di3"}},
                                data=dx)
result10 = model10.fit()
print(result10.summary())

model11 = sm.MixedLM.from_formula("bp ~ RIDAGEYR + female + C(bpt) + BMXBMI",
                                groups="SEQN", re_formula="1",
                                vc_formula={"sy": "0+sy", "di": "0+di",
                                             "sy": "0+sy1+sy2+sy3"}},
                                data=dx)
result11 = model11.fit()
print(result11.summary())

```

Mixed Linear Model Regression Results

```

=====
Model:                MixedLM Dependent Variable: bp
No. Observations: 5000   Method:                REML
No. Groups:           837   Scale:                16.1732
Min. group size:      2    Likelihood:            -17769.2792
Max. group size:      6    Converged:              Yes
Mean group size:      6.0

```

```

-----
                Coef.  Std.Err.   z    P>|z|  [0.025  0.975]
-----
Intercept      46.301    1.616  28.655  0.000  43.134  49.468
C(bpt)[T.SY]   52.003    0.599  86.752  0.000  50.828  53.178
RIDAGEYR        0.289    0.021  13.991  0.000   0.248   0.329
female        -2.153    0.825  -2.611  0.009  -3.769  -0.537
BMXBMI          0.369    0.061   6.082  0.000   0.250   0.488
SEQN Var       63.534    1.996
di Var         134.571    2.807
dye Var         21.006    0.454
sy Var         148.524    2.930

```

```

/nfs/kshedden/python3/lib/python3.7/site-packages/statsmodels/base/model.py:512:
ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check
mle_retvals

```

```

    "Check mle_retvals", ConvergenceWarning)

```

```

/nfs/kshedden/python3/lib/python3.7/site-
packages/statsmodels/regression/mixed_linear_model.py:2059: ConvergenceWarning:
Retrying MixedLM optimization with lbfgs

```

```

    ConvergenceWarning)

```

```

/nfs/kshedden/python3/lib/python3.7/site-packages/statsmodels/base/model.py:512:
ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check
mle_retvals

```

```

    "Check mle_retvals", ConvergenceWarning)

```

```

/nfs/kshedden/python3/lib/python3.7/site-
packages/statsmodels/regression/mixed_linear_model.py:2059: ConvergenceWarning:
Retrying MixedLM optimization with cg
    ConvergenceWarning)
/nfs/kshedden/python3/lib/python3.7/site-packages/statsmodels/base/model.py:512:
ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check
mle_retvals
    "Check mle_retvals", ConvergenceWarning)
/nfs/kshedden/python3/lib/python3.7/site-
packages/statsmodels/regression/mixed_linear_model.py:2063: ConvergenceWarning:
MixedLM optimization failed, trying a different optimizer may help.
    warnings.warn(msg, ConvergenceWarning)
/nfs/kshedden/python3/lib/python3.7/site-
packages/statsmodels/regression/mixed_linear_model.py:2075: ConvergenceWarning:
Gradient optimization failed, |grad| = 145.966482
    warnings.warn(msg, ConvergenceWarning)
/nfs/kshedden/python3/lib/python3.7/site-
packages/statsmodels/regression/mixed_linear_model.py:2115: ConvergenceWarning:
The Hessian matrix at the estimated parameter values is not positive definite.
    warnings.warn(msg, ConvergenceWarning)

```

Mixed Linear Model Regression Results

```

=====
Model:                MixedLM Dependent Variable: bp
No. Observations: 5000    Method:                REML
No. Groups:            837    Scale:                28.7455
Min. group size:      2      Likelihood:            -17944.9876
Max. group size:      6      Converged:              No
Mean group size:      6.0

-----
               Coef.   Std.Err.    z    P>|z| [0.025 0.975]
-----
Intercept      46.045     1.617   28.471 0.000  42.876  49.215
C(bpt)[T.SY]   52.017     0.518  100.461 0.000  51.003  53.032
RIDAGEYR        0.302     0.021   14.605 0.000   0.261   0.342
female        -2.216     0.827   -2.681 0.007  -3.836  -0.596
BMXBMI         0.360     0.061    5.918 0.000   0.241   0.479
SEQN Var       83.637     1.215
di Var        112.382     1.240
sy Var         92.482     1.220
sy Var         0.641
=====

```

```

/nfs/kshedden/python3/lib/python3.7/site-
packages/statsmodels/base/model.py:1286: RuntimeWarning: invalid value
encountered in sqrt
    bse_ = np.sqrt(np.diag(self.cov_params()))

```